

The Council of Nark

Why try several specialist reviewers instead of one general review?

The idea grew out of Fable 5's multi-agent review mode. It ran while I planned across projects on paper.



The workflow already mixed reviewers

The proposal made an informal practice explicit.

01

Fable 5

Several agent reviews, followed by one combined response.

02

Gemini

A separate model for research, challenge and a second opinion.

03

HK-47

A fixed role that looks for code and process we do not need.

PROPOSAL

Give each reviewer one clear responsibility, run relevant reviewers independently, then combine their findings once.

Meet the council



The roles make the review lenses easy to recall

Seven memorable characters. Seven fixed review responsibilities.

Who reviews what?

Each reviewer has a narrow lens. None of them edits the code, plan or document.

	REVIEWER	RESPONSIBILITY
	HK-47	Remove unnecessary complexity and speculative machinery.
	K-2SO	Check correctness, tests, observability and data edge cases.
	GLaDOS	Check architecture, contracts, coupling and migrations.
	C-3PO	Check secrets, personal data, permissions and destructive actions.
	Bender	Check compute use and avoidable cost.
	Holly	Check what breaks after months of neglect or staff change.
	WALL-E	Check durable technical language and translate the final review.

Holly and WALL-E work at different points



Holly

INDEPENDENT PANEL REVIEW

What happens after six months of ordinary neglect?

- Forgotten toggles and manual steps
- State drift and missing ownership
- Noisy alerts and difficult handovers



WALL-E

WRITING REVIEW AND FINAL TRANSLATION

Can a tired reader understand the decision once?

- Unclear conditions and vague warnings
- Dense jargon and inconsistent terms
- A plain-language version of the combined review

independent findings



arbiter merges and ranks




WALL-E translates



human decision

One review can produce different outcomes

VISUAL MODEL OF THE ORIGINAL HYPOTHESIS




The same model and prompt can give a stronger or weaker review on different runs.

The original shorthand was a **cone of possible outcomes**.

Better instructions can move the whole range, but they do not remove variation.

A serial chain can carry mistakes forward

ORIGINAL HYPOTHESIS: SERIAL REVIEW




In a chain, each reviewer sees the previous answer as part of its starting point.

A later reviewer can correct an error, repeat it or remove a valid finding.

The order of reviewers can therefore change the final review.

One arbiter combines independent reviews

ORIGINAL HYPOTHESIS: INDEPENDENT REVIEW AND FUSION



Every specialist receives the same source, without seeing another review.


The arbiter removes duplicates and keeps supported findings raised by only one reviewer.

Hypothesis: this raises useful coverage and reduces weak final reviews.

Flower, not a chain


The same reviewers can be connected in two different ways.

FAN OUT, THEN FUSE



Each reviewer sees the source. The arbiter compares independent findings once.

PASS THE ANSWER ALONG



Each reviewer inherits earlier output. Sequence and omissions can affect the result.

A testable proposal, not a conclusion

01

Specialist roles

Compare them with the same number of repeated general reviews.

02

Character wording

Keep the functional instructions fixed and change only the character text.

03

Review structure

Compare fan-out and fusion with every order of a short chain.

Keep extra reviewers only when they improve useful findings enough to justify their time and cost.

