

02 / THE FAIR TEST

Put the Council on Trial

Separate more effort from better review.

Change one thing at a time

Question

Comparison

Does repetition help?

One omnibus vs repeated omnibus

Do roles help?

Repeated omnibus vs specialists

Do characters help?

Same kernels, different wrappers

Does fusion help?

Raw findings vs final verdict

A score needs an answer key

- Three frozen synthetic review packets.
- Planted defects and deliberately clean facts.
- Supported findings earn credit; invented faults cost precision.

Repeated calls measure sampling variation, not new tasks.

Equal scores can hide different reviews

F1 balances supported findings and missed defects.

- It does not tell us which defects were found.
- It does not measure whether a remedy is good.
- It does not price the reviewer's extra effort.

The measurement needs its own controls

- Freeze inputs; preserve every attempt.
- Keep answer keys outside the model's reach.
- Lock independent ratings before revealing conditions.

Current published scores use exploratory LLM ratings.

“No clear winner” is not equivalence

Define a meaningful difference before collecting results.

| An interval must be narrow enough to rule out a useful gain or loss.