

03 / THE EVIDENCE

The Experiments Failed. A Lot.

What broke, what survived, and what changed.

First, make sure an answer is an answer

Failure

Schema rejected before inference

Prompt echo looked like output

Duplicate complaints inflated noise

Course correction

Check the adapter first

Parse assistant events only

Cluster unsupported claims

The partial 71/81 run was sealed and excluded.

Eight calls matched one review's score

0.860

One functional review

0.860

Seven reviews + fusion

| 8.9× the recorded cost in the Haiku smoke.

Three tasks, one draw. Equal F1 did not mean identical findings.

The next smoke favoured specialists

Gemma final review	Mean F1
One functional omnibus	0.804
Repeated omnibus + fusion	0.863
Functional specialists + fusion	0.911
Fictional specialists + fusion	0.857

One draw on three tasks. The ordering changed; proof did not arrive.

K-2SO lost the repeated comparison

0.790

Functional correctness

0.748

Same kernel + K-2SO

| 2 fictional wins · 16 ties · 12 functional wins

30 pairs, but only three tasks. One model and one persona.

Better isolation enabled the next test

- History-free source, with upstream answers removed.
- Bounded source access and offline regression tests.
- Gortex's parent fails; its evidence revision passes.

| Zero ecological respondent calls so far.