

04 / THE PIVOT

What Would Settle It?

The next comparison, defined before the next call.

Give one agent a fair chance

Proposed arm

What it isolates

One functional review

Cheap baseline

One agent, self-review

Fewer-agent alternative

Three omnibus + fusion

Independent sampling

Three specialists + fusion

Functional roles

Same specialists + personas

Personality package

Real projects make the test relevant

Candidate project

Frozen task

Manim

Animation ordering regression

dlt

Paginator stopping precedence

Gortex

Unicode tokenizer panic

turbovec

Stale temporary-file cleanup

Candidates are not outcomes. Each needs a validated test environment.

Define what would change the decision

- Better: a useful quality gain with acceptable failure risk.
- Equivalent: uncertainty fits inside a declared small margin.
- Cheaper substitute: quality holds up and resources fall.

| Apply the rule after the confirmatory run.

Run feasibility, then confirmation

4×2

Tasks $\times$ repeats for feasibility

57^*

Illustrative independent-task target

*Freeze the final count from the declared precision rule.

Next: run the ecological comparison

Finish the runner.

Freeze the tasks.

Collect blinded human ratings.

| Then publish Part 5: the result.