

04 / THE CONCLUSION

What Would Settle It?

A smaller final test. An honest stopping rule.

Give one agent a fair chance

Proposed arm

What it isolates

One functional review

Cheap baseline

One agent, self-review

Fewer-agent alternative

Three omnibus + fusion

Independent sampling

Three specialists + fusion

Functional roles

Same specialists + personas

Personality package

Real projects make the test relevant

Candidate project

Frozen task

Manim

Animation ordering regression

dlt

Paginator stopping precedence

Gortex

Unicode tokenizer panic

turbovec

Stale temporary-file cleanup

Candidates are not outcomes. Each needs a validated test environment.

Define what would change the decision

- Better: a useful quality gain with acceptable failure risk.
- Equivalent: uncertainty fits inside a declared small margin.
- Cheaper substitute: quality holds up and resources fall.

Otherwise, publish “inconclusive”.

A small pilot cannot prove a small gap

4 × 2

Tasks × repeats for feasibility

136

Planned stages, not API calls

Confirmation needs a sample sized for precision, then one final analysis.

The council has not earned a default

Start with one functional review.

Escalate when the task warrants it.

A practical choice under uncertainty.

Not a claim of proven equivalence.